

# Analysis of activations in neural networks

Jan Ruman

30.8.2021

## Abstract

In this work, we attempt to analyze the properties of activations inside of a neural network. Embeddings of activations inside trained and untrained neural network are compared across different layers. According to our results, activations from the trained network systematically outperform activations from the untrained network in terms of embedding quality.

## 1 Introduction

A neural network applies a sequence of non-linear transformations to its input in order to produce an output. This procedure can be (and is) done for all examples in the dataset. By storing the intermediate transformations, we end up with multiple different representations of the original dataset (possibly compressed and/or missing some information).

This work aims to analyze these representations in terms of their structure. We are primarily interested in the way examples with the same class are distributed when using the new representations.

## 2 Methodology

We will compare the data representations across two dimensions: 1) the depth of layer from which the representations were acquired and 2) whether the network was or was not trained. This means we will need a neural network to be trained and datasets to train it on.

The representations are going to have many dimensions - it will be more practical to use dimensionality reduction techniques that retain the structure of the data. For that purpose, we are going to use t-SNE and UMAP algorithms.

We need some way of comparing the structure of different representations. Plotting embeddings acquired using dimensionality reduction techniques is one of them, but some more rigorous approaches will be necessary to acquire more reliable results. In order to measure the quality of a representation, i.e., how well it clusters the same examples together and keeps different examples away, we introduce several metrics in the next section.

A convolutional neural network with ReLU activation function will be used. The first layer of the network is a convolutional layer and it consists of 4 filters of size 5. The second layer is a subsampling layer with kernel size 2. The third layer is convolutional with 16 filters of size 3. The fourth layer is subsampling layer with kernel size 2. The fifth layer is convolutional with 64 filters of size 5. The sixth layer is fully connected. ReLU activation function is applied after second, fourth, fifth and sixth layer.

As representations, we will use:

- the original dataset
- feature maps of the first layer (one for each filter)
- feature maps of the second layer (one for each filter)
- activations in the third layer (consisting of 64 1x1 feature maps)

We will use MNIST and FashionMNIST datasets. Because of memory limitations, we will only work with a subset of 5 000 examples of each dataset.

### 3 Embedding quality

To measure the quality of a representation, we introduce several metrics that capture the representation's properties. These properties then indicate how good the representation is.

All except the last metric have distance as a unit. This is potentially problematic as the absolute distance between points tells us nothing if we do not know the distances between other points. It is thus necessary to divide each of these metrics by the average distance between two points; this will ensure that even though two embeddings have different scales, we can still compare their metrics.

#### 3.1 Distances between points

The first two metrics are concerned with an average distance between two points. An average distance between points of the same class ( $d_s$ ) and an average distance between points of different classes ( $d_d$ ) are measured. It is quite obvious which one we would like to minimize and which one to maximize.

#### 3.2 Cluster centroids

We can look at points belonging to the same class as a cluster. We introduce two metrics using centroids of these clusters: 1) the total distance from points of a class to its centroid ( $c_s$ ) and 2) the total distance between different centroids ( $c_d$ ).

### 3.3 k-nearest neighbors

We can use the k-nearest neighbors algorithm on the embedding to classify examples in the dataset. The accuracy of k-nn ( $acc_{knn}$ ) should correlate with the quality of the embedding.

## 4 Results

Results presented in this part were acquired by running the `scripts/activations_cnn.py` script. Output of this script used in this document is part of the repository (in particular, the output is stored in the `out/activations_cnn/` directory); rerunning the script might give slightly different results.

In this work we will present only fraction of the visualizations actually computed because of lack of space. The reader might see the rest of the visualizations in the output directory mentioned above.

### 4.1 MNIST

First, we will look at embeddings produced by t-SNE. As there are 22 embeddings in total, it wouldn't be practical to show them all. Instead, we are going to show one embedding per layer of the neural network. These can be seen in Figures 1 to 7.

We might make an interesting observation - in the trained neural network, the representations don't seem to get simply better the deeper the layer is. Representations in the third layer seem to be the best (in terms of distances between examples of same / different class), but representations in the second and third layer seem to be worse than representations for the original dataset. This could be explained by the fact that a lot of information can be lost by applying a filter to an image; combining the outputs of the filters should then result in greater embedding quality. This would suggest that computing an average of values of individual metrics over the whole layer might not be the best way to actually evaluate the quality of embeddings for the whole layer; how that could be done is out of the scope of this work.

The representations obtained from the untrained network follow the same pattern with two important distinctions:

1. The representations are generally better in the trained network.
2. The representation quality is worse in the third layer of an untrained neural network than it is in the case of the original dataset; the opposite holds in the case of the trained network.

Plotting the embeddings gives us a rough idea of the quality of the embeddings. However, we would also like to see values of metrics introduced in the previous section. Figures 8 to 11 present tables with values of those metrics. There are two values per layer and metric - the first one corresponds to the




Figure 1: Embeddings of the original dataset (t-SNE).




Figure 2: Embeddings of activations in the first layer of a neural network trained on MNIST (t-SNE).




Figure 3: Embeddings of activations in the second layer of a neural network trained on MNIST (t-SNE).




Figure 4: Embeddings of activations in the third layer of a neural network trained on MNIST (t-SNE).




Figure 5: Embeddings of activations in the first layer of an untrained neural network (t-SNE).




Figure 6: Embeddings of activations in the second layer of an untrained neural network (t-SNE).




Figure 7: Embeddings of activations in the third layer of an untrained neural network (t-SNE).

	ds	dd	cs	cd	acc_knn
filter 1	0.3773 / 0.3773	1.0695 / 1.0695	0.2664 / 0.2664	0.9748 / 0.9748	0.9398 / 0.9398

Figure 8: Values of metrics for embeddings acquired by t-SNE, original dataset.

trained network, the second one to the untrained network. As it is not very easy to see what is the general trend in these values, we also provide plots with average values of metrics in each layer. We only provide plots for the  $d_s$  and  $d_d$  metrics for lack of space; however, the same trend can be observed in the rest of the metrics too. Figures 12 and 13 contain these plots.

The trend we are seeing both from the tables and from the plots is the same one we noted when inspecting the embeddings visually.

Even the representations in the third layer of the trained neural network seem to contain some outliers - points that are in a wrong cluster. For example, we might notice that there are several examples with true label 3 in the cluster of examples labeled 7. One such example can be seen in Figure 14.

The figure also contains two other examples, which show what an average example of 3 and 7 looks like. What we might notice is that the outlier does seem to contain features characteristic for both 3 and 7 and it might not be immediately clear, even to a human, what number it really is.

It might also be of interest to us what do the individual filters do. Figures 15 to 17 visualize the activations of the network for an example from the dataset.

	ds	dd	cs	cd	acc_knn
filter 1	0.3550 / 0.4218	1.0720 / 1.0645	0.2589 / 0.3085	0.9758 / 0.9486	0.9340 / 0.9328
filter 2	0.3771 / 0.3643	1.0695 / 1.0709	0.2714 / 0.2612	0.9805 / 0.9761	0.9472 / 0.9428
filter 3	0.3316 / 0.3296	1.0746 / 1.0748	0.2351 / 0.2346	0.9958 / 0.9952	0.9394 / 0.9448
filter 4	0.4326 / 0.4920	1.0633 / 1.0567	0.3155 / 0.3670	0.9331 / 0.8796	0.9214 / 0.8996

Figure 9: Values of metrics for embeddings acquired by t-SNE, first layer.

	ds	dd	cs	cd	acc_knn
filter 1	0.4455 / 0.5469	1.0619 / 1.0506	0.3277 / 0.4174	0.9137 / 0.7944	0.8794 / 0.8224
filter 2	0.5410 / 0.6565	1.0512 / 1.0383	0.4017 / 0.4807	0.8445 / 0.7281	0.8516 / 0.7182
filter 3	0.5855 / 0.5865	1.0463 / 1.0461	0.4330 / 0.4219	0.8116 / 0.8147	0.8442 / 0.8266
filter 4	0.5770 / 0.4310	1.0472 / 1.0635	0.4264 / 0.3080	0.8110 / 0.9249	0.8512 / 0.8782
filter 5	0.4406 / 0.5063	1.0624 / 1.0551	0.3176 / 0.3809	0.9215 / 0.8540	0.8920 / 0.8746
filter 6	0.4279 / 0.4685	1.0638 / 1.0593	0.3054 / 0.3410	0.9402 / 0.9229	0.8794 / 0.8508
filter 7	0.4393 / 0.4066	1.0626 / 1.0662	0.3081 / 0.2887	0.9218 / 0.9443	0.8798 / 0.8950
filter 8	0.4219 / 0.5111	1.0645 / 1.0546	0.3107 / 0.3710	0.9338 / 0.8676	0.8882 / 0.7876
filter 9	0.4777 / 0.4045	1.0583 / 1.0664	0.3469 / 0.2893	0.8954 / 0.9550	0.8532 / 0.8878
filter 10	0.4539 / 0.8760	1.0609 / 1.0138	0.3256 / 0.6570	0.9248 / 0.4283	0.8688 / 0.4660
filter 11	0.5357 / 0.7257	1.0518 / 1.0306	0.3995 / 0.5299	0.8440 / 0.6375	0.8642 / 0.6910
filter 12	0.4215 / 0.7265	1.0646 / 1.0305	0.3121 / 0.5783	0.9274 / 0.6335	0.8992 / 0.4774
filter 13	0.4534 / 0.6405	1.0610 / 1.0401	0.3194 / 0.4892	0.9305 / 0.7075	0.8568 / 0.8198
filter 14	0.6861 / 0.5102	1.0350 / 1.0547	0.5150 / 0.3752	0.6828 / 0.8604	0.7858 / 0.8278
filter 15	0.5404 / 0.6304	1.0513 / 1.0412	0.3984 / 0.4869	0.8416 / 0.7232	0.8630 / 0.8064
filter 16	0.4875 / 0.5654	1.0572 / 1.0485	0.3595 / 0.4239	0.8901 / 0.8020	0.8736 / 0.8124

Figure 10: Values of metrics for embeddings acquired by t-SNE, second layer.

	ds	dd	cs	cd	acc_knn
filter 1	0.2370 / 0.5046	1.0851 / 1.0553	0.1678 / 0.3772	1.0544 / 0.8534	0.9854 / 0.8654

Figure 11: Values of metrics for embeddings acquired by t-SNE, third layer.

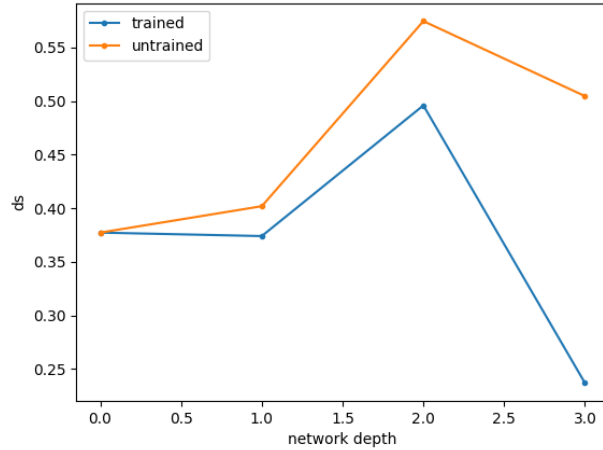


Figure 12: Average value of the  $d_s$  metric (t-SNE).






Figure 13: Average value of the  $d_d$  metric (t-SNE).




Figure 14: From left: an example of 3, an example of 7, the missclassified example.




Figure 15: An example of 5.




Figure 16: Activations in the first layer of the network for the example.




Figure 17: Activations in the second layer of the network for the example.




Figure 18: Embeddings of the original dataset (UMAP).

In the first layer, it seems that the first and third filter detects some variation of a horizontal line, and it also seems that the second and fourth filter detects some variation of a diagonal line. The activations in the second layer are not as easy to analyze and it is not clear what purpose each filter serves.

Using UMAP to acquire embeddings, we get very similar results. Figures 18 to 24 show the computed embeddings.

Again, we would like to confirm our observations by comparing values of metrics over the computed embeddings. These can be seen in Figures 25 to 30.

Again, the trend we are seeing both from the tables and from the plots is the same one we noted when inspecting the embeddings visually.

## 4.2 FashionMNIST

Again, we start by looking at embeddings produced by t-SNE. A single embedding is shown per layer of the network in Figures 31 to 37.

The trend we observed in the case of MNIST dataset is present in this case too. To make sure this is really the case, we compute values of metrics of embedding quality and compare between the trained and the untrained network. See Figures 38 to 43.

We would like to look again a bit more closely at the outliers present in the embeddings acquired from the third layer of the trained network. We notice there is a big overlap between the Shirt and T-shirt/top clusters. Figure 44




Figure 19: Embeddings of activations in the first layer of a neural network trained on MNIST (UMAP).




Figure 20: Embeddings of activations in the second layer of a neural network trained on MNIST (UMAP).

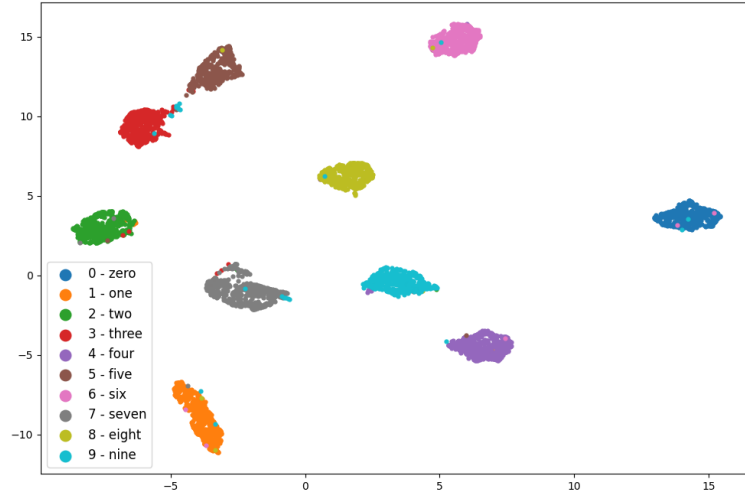


Figure 21: Embeddings of activations in the third layer of a neural network trained on MNIST (UMAP).




Figure 22: Embeddings of activations in the first layer of an untrained neural network (UMAP).




Figure 23: Embeddings of activations in the second layer of an untrained neural network (UMAP).




Figure 24: Embeddings of activations in the third layer of an untrained neural network (UMAP).



	ds	dd	cs	cd	acc_knn
filter 1	0.2759 / 0.2721	1.0808 / 1.0812	0.1921 / 0.1894	1.0125 / 1.0146	0.9268 / 0.9396

Figure 25: Values of metrics for embeddings acquired by UMAP, original dataset.

	ds	dd	cs	cd	acc_knn
filter 1	0.2682 / 0.2983	1.0817 / 1.0783	0.1902 / 0.2089	1.0118 / 1.0004	0.9278 / 0.9028
filter 2	0.2397 / 0.2550	1.0848 / 1.0831	0.1682 / 0.1776	1.0326 / 1.0227	0.9492 / 0.9434
filter 3	0.2601 / 0.2325	1.0826 / 1.0856	0.1814 / 0.1625	1.0226 / 1.0339	0.9296 / 0.9350
filter 4	0.3193 / 0.3608	1.0760 / 1.0713	0.2261 / 0.2589	1.0041 / 0.9694	0.8956 / 0.8800

Figure 26: Values of metrics for embeddings acquired by UMAP, first layer.

	ds	dd	cs	cd	acc_knn
filter 1	0.3678 / 0.4829	1.0705 / 1.0577	0.2664 / 0.3686	0.9489 / 0.8499	0.8736 / 0.7978
filter 2	0.5066 / 0.6238	1.0551 / 1.0420	0.3776 / 0.4587	0.8616 / 0.7358	0.8182 / 0.6976
filter 3	0.4963 / 0.5375	1.0562 / 1.0516	0.3586 / 0.3827	0.8764 / 0.8453	0.8196 / 0.7820
filter 4	0.5073 / 0.3684	1.0550 / 1.0705	0.3638 / 0.2636	0.8848 / 0.9623	0.7918 / 0.8422
filter 5	0.3663 / 0.4185	1.0707 / 1.0649	0.2597 / 0.3025	0.9701 / 0.9230	0.8814 / 0.8440
filter 6	0.3964 / 0.4408	1.0674 / 1.0624	0.2839 / 0.3222	0.9544 / 0.9276	0.8694 / 0.8250
filter 7	0.4176 / 0.3447	1.0650 / 1.0731	0.2937 / 0.2397	0.9295 / 0.9661	0.8600 / 0.8856
filter 8	0.3533 / 0.4895	1.0722 / 1.0570	0.2497 / 0.3506	0.9644 / 0.8867	0.8816 / 0.7530
filter 9	0.4221 / 0.3705	1.0645 / 1.0702	0.3009 / 0.2604	0.9351 / 0.9663	0.8312 / 0.8516
filter 10	0.4524 / 0.8812	1.0611 / 1.0133	0.3244 / 0.6720	0.9181 / 0.4235	0.8534 / 0.4684
filter 11	0.4570 / 0.6986	1.0606 / 1.0336	0.3371 / 0.5102	0.9120 / 0.6830	0.8444 / 0.6626
filter 12	0.3639 / 0.7084	1.0710 / 1.0325	0.2654 / 0.5821	0.9553 / 0.6493	0.8930 / 0.4898
filter 13	0.4148 / 0.5297	1.0653 / 1.0525	0.2965 / 0.3886	0.9470 / 0.8504	0.8350 / 0.7884
filter 14	0.6577 / 0.4621	1.0382 / 1.0600	0.4898 / 0.3369	0.7099 / 0.8882	0.7464 / 0.7920
filter 15	0.4817 / 0.5614	1.0578 / 1.0489	0.3511 / 0.4265	0.8936 / 0.8036	0.8394 / 0.7754
filter 16	0.4008 / 0.4948	1.0669 / 1.0564	0.2897 / 0.3583	0.9410 / 0.8692	0.8596 / 0.7852

Figure 27: Values of metrics for embeddings acquired by UMAP, second layer.

	ds	dd	cs	cd	acc_knn
filter 1	0.1172 / 0.4456	1.0985 / 1.0619	0.0799 / 0.3288	1.0819 / 0.8912	0.9834 / 0.8310

Figure 28: Values of metrics for embeddings acquired by UMAP, third layer.




Figure 29: Average value of the  $d_s$  metric (UMAP).

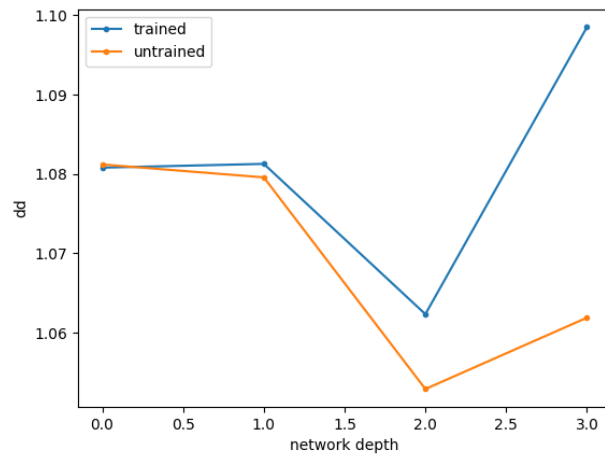


Figure 30: Average value of the  $d_d$  metric (UMAP).




Figure 31: Embeddings of the original dataset (t-SNE).




Figure 32: Embeddings of activations in the first layer of a neural network trained on FashionMNIST (t-SNE).




Figure 33: Embeddings of activations in the second layer of a neural network trained on FashionMNIST (t-SNE).




Figure 34: Embeddings of activations in the third layer of a neural network trained on FashionMNIST (t-SNE).




Figure 35: Embeddings of activations in the first layer of an untrained neural network (t-SNE).




Figure 36: Embeddings of activations in the second layer of an untrained neural network (t-SNE).




Figure 37: Embeddings of activations in the third layer of an untrained neural network (t-SNE).

	ds	dd	cs	cd	acc_knn
filter 1	0.4190 / 0.4190	1.0646 / 1.0646	0.3117 / 0.3117	0.9527 / 0.9527	0.8328 / 0.8328

Figure 38: Values of metrics for embeddings acquired by t-SNE, original dataset.

	ds	dd	cs	cd	acc_knn
filter 1	0.9873 / 0.9848	1.0014 / 1.0017	0.7952 / 0.8904	0.3079 / 0.3053	0.1012 / 0.1084
filter 2	0.4394 / 0.4577	1.0623 / 1.0603	0.3246 / 0.3335	0.9408 / 0.9308	0.8158 / 0.7830
filter 3	0.9009 / 0.7268	1.0110 / 1.0304	0.7932 / 0.5390	0.4736 / 0.6707	0.2068 / 0.5576
filter 4	0.3994 / 0.4192	1.0668 / 1.0646	0.2935 / 0.3081	0.9666 / 0.9551	0.8340 / 0.8134

Figure 39: Values of metrics for embeddings acquired by t-SNE, first layer.

	ds	dd	cs	cd	acc_knn
filter 1	0.5084 / 0.4670	1.0546 / 1.0593	0.3740 / 0.3380	0.8808 / 0.9186	0.7244 / 0.7672
filter 2	0.5349 / 0.5292	1.0517 / 1.0523	0.3940 / 0.3880	0.8599 / 0.8772	0.7172 / 0.7110
filter 3	0.5359 / 0.9431	1.0516 / 1.0063	0.3937 / 0.6762	0.8661 / 0.2595	0.7244 / 0.1690
filter 4	0.5437 / 1.0017	1.0507 / 0.9998	0.4131 / 0.9410	0.8399 / 0.2029	0.7132 / 0.1016
filter 5	0.4465 / 0.4365	1.0615 / 1.0626	0.3263 / 0.3177	0.9430 / 0.9420	0.7810 / 0.7814
filter 6	0.5255 / 0.4312	1.0527 / 1.0632	0.3843 / 0.3120	0.8669 / 0.9430	0.7216 / 0.7812
filter 7	0.5510 / 0.5872	1.0499 / 1.0459	0.4042 / 0.4306	0.8412 / 0.8213	0.6898 / 0.6542
filter 8	0.5884 / 0.7899	1.0458 / 1.0234	0.4321 / 0.6837	0.8244 / 0.6114	0.7060 / 0.2494
filter 9	0.5553 / 0.9931	1.0494 / 1.0008	0.4111 / 0.8814	0.8531 / 0.2709	0.6984 / 0.1060
filter 10	0.5370 / 0.4733	1.0515 / 1.0585	0.3956 / 0.3455	0.8394 / 0.9102	0.7230 / 0.7714
filter 11	0.5395 / 0.5819	1.0512 / 1.0465	0.3998 / 0.4373	0.8721 / 0.8305	0.7276 / 0.7010
filter 12	0.6236 / 0.9583	1.0418 / 1.0046	0.4656 / 0.7442	0.7961 / 0.3636	0.6950 / 0.1350
filter 13	0.6807 / 0.6936	1.0355 / 1.0341	0.5104 / 0.5441	0.6876 / 0.6840	0.6796 / 0.4448
filter 14	0.5985 / 0.9720	1.0446 / 1.0031	0.4420 / 0.8357	0.8050 / 0.3537	0.6876 / 0.1310
filter 15	0.4729 / 0.4339	1.0586 / 1.0629	0.3484 / 0.3145	0.9214 / 0.9404	0.7624 / 0.7824
filter 16	0.6315 / 0.7519	1.0410 / 1.0276	0.4701 / 0.5599	0.7535 / 0.6404	0.7232 / 0.5594

Figure 40: Values of metrics for embeddings acquired by t-SNE, second layer.

	ds	dd	cs	cd	acc_knn
filter 1	0.3295 / 0.4415	1.0745 / 1.0621	0.2402 / 0.3221	1.0055 / 0.9358	0.8874 / 0.7856

Figure 41: Values of metrics for embeddings acquired by t-SNE, third layer.




Figure 42: Average value of the  $d_s$  metric (t-SNE).




Figure 43: Average value of the  $d_d$  metric (t-SNE).

shows examples of these classes (as well as an example that would be missclassified as Shirt even though it is a T-shirt/top).

We would also like to see what do the individual feature maps look like for some example. We will use a T-shirt/top as such an example. Figures 45 to 47 contain these visualizations.

Regarding the first layer, it is not clear what does the first filter detect; the second filter seems to detect a transition from light pixels to dark pixels in the top-bottom direction, and the third and fourth filter seem to detect bottom right and upper left edge of a group of light pixels, respectively. In the second layer, it is very hard to tell what meaning each of the filters has.

Using UMAP to acquire embeddings, we get very similar results. Figures 48




Figure 44: From left: an example of T-shirt/top, an example of Shirt, the missclassified example.






Figure 45: An example of a T-shirt/top.




Figure 46: Activations in the first layer of the network for the example.




Figure 47: Activations in the second layer of the network for the example.

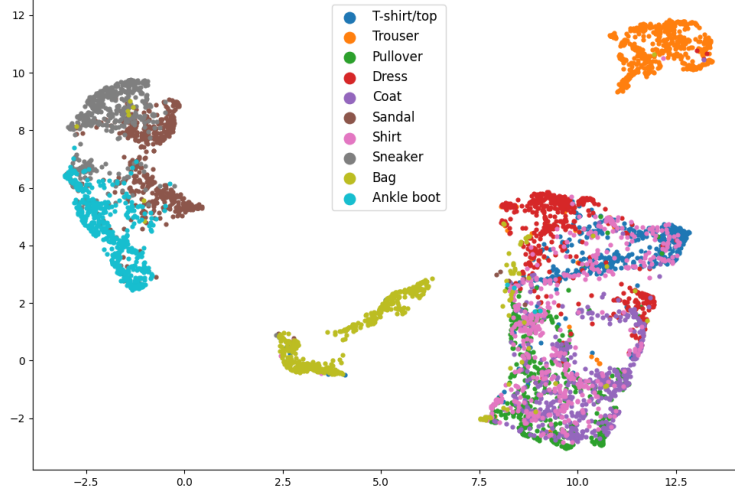


Figure 48: Embeddings of the original dataset (UMAP).

to 54 show the computed embeddings.

Again, we would like to confirm our observations by comparing values of metrics over the computed embeddings. These can be seen in Figures 54 to 59.

Again, the trend we are seeing both from the tables and from the plots is the same one we noted when inspecting the embeddings visually.

## 5 Conclusions

In this work, we attempted to tackle the problem of data representation using activations of a neural network. To find out whether these representations are superior to their original counterparts, we compared them both visually and using several metrics of embedding quality.

According to our results, the quality of embeddings acquired from a trained network is in general better than the quality of embeddings acquired from an untrained network irrespectively of depth of layer. Moreover, the quality of embeddings acquired from the last layer of a network is generally better than the embeddings computed from the original dataset in the case of trained neural network, but not in the case of the untrained one.

Our results also suggest that by applying a convolutional filter a feature map missing part of the information present in the original image is obtained, and both in the case of trained and untrained network different filters seem to extract different information from the original image.




Figure 49: Embeddings of activations in the first layer of a neural network trained on FashionMNIST (UMAP).




Figure 50: Embeddings of activations in the second layer of a neural network trained on FashionMNIST (UMAP).




Figure 51: Embeddings of activations in the third layer of a neural network trained on FashionMNIST (UMAP).




Figure 52: Embeddings of activations in the first layer of an untrained neural network (UMAP).




Figure 53: Embeddings of activations in the second layer of an untrained neural network (UMAP).




Figure 54: Embeddings of activations in the third layer of an untrained neural network (UMAP).

	ds	dd	cs	cd	acc_knn
filter 1	0.2715 / 0.2744	1.0810 / 1.0807	0.1991 / 0.2017	1.0162 / 1.0154	0.8014 / 0.8078

Figure 55: Values of metrics for embeddings acquired by UMAP, original dataset.

	ds	dd	cs	cd	acc_knn
filter 1	1.0004 / 0.9999	1.0000 / 1.0000	0.7222 / 0.7213	0.0382 / 0.0504	0.2832 / 0.2762
filter 2	0.2912 / 0.3710	1.0788 / 1.0699	0.2150 / 0.2711	1.0145 / 0.9738	0.7710 / 0.7546
filter 3	0.9491 / 0.8821	1.0057 / 1.0131	0.6951 / 0.6522	0.2907 / 0.2972	0.3304 / 0.5526
filter 4	0.2548 / 0.2902	1.0828 / 1.0789	0.1850 / 0.2116	1.0243 / 1.0118	0.8048 / 0.7932

Figure 56: Values of metrics for embeddings acquired by UMAP, first layer.

	ds	dd	cs	cd	acc_knn
filter 1	0.3698 / 0.4052	1.0701 / 1.0661	0.2717 / 0.2918	0.9546 / 0.9488	0.7144 / 0.7578
filter 2	0.4812 / 0.4893	1.0577 / 1.0568	0.3524 / 0.3536	0.8889 / 0.8878	0.6960 / 0.6816
filter 3	0.5024 / 0.9913	1.0553 / 1.0010	0.3714 / 0.7064	0.8888 / 0.1138	0.6978 / 0.2986
filter 4	0.4162 / 1.0003	1.0649 / 1.0000	0.2992 / 0.7222	0.9375 / 0.0323	0.6936 / 0.2690
filter 5	0.3514 / 0.3352	1.0721 / 1.0739	0.2563 / 0.2453	0.9913 / 0.9906	0.7658 / 0.7568
filter 6	0.4769 / 0.3038	1.0581 / 1.0774	0.3494 / 0.2184	0.8998 / 1.0044	0.6982 / 0.7624
filter 7	0.4869 / 0.5378	1.0570 / 1.0514	0.3545 / 0.3946	0.8674 / 0.8323	0.6756 / 0.6348
filter 8	0.5598 / 0.8257	1.0489 / 1.0194	0.4134 / 0.6523	0.8421 / 0.5345	0.6744 / 0.3618
filter 9	0.4940 / 0.9999	1.0562 / 1.0000	0.3639 / 0.7214	0.8914 / 0.0486	0.6700 / 0.2740
filter 10	0.4605 / 0.3935	1.0600 / 1.0674	0.3337 / 0.2856	0.8934 / 0.9619	0.7012 / 0.7526
filter 11	0.4727 / 0.5849	1.0586 / 1.0461	0.3525 / 0.4383	0.9189 / 0.8520	0.7118 / 0.6840
filter 12	0.6188 / 0.9937	1.0424 / 1.0007	0.4627 / 0.7066	0.7985 / 0.1411	0.6734 / 0.2878
filter 13	0.6530 / 0.6921	1.0386 / 1.0342	0.4908 / 0.5553	0.7295 / 0.6328	0.6484 / 0.5086
filter 14	0.5699 / 0.9997	1.0478 / 1.0000	0.4244 / 0.7212	0.8216 / 0.0562	0.6620 / 0.2810
filter 15	0.4351 / 0.3349	1.0628 / 1.0739	0.3201 / 0.2423	0.9299 / 0.9879	0.7588 / 0.7664
filter 16	0.6159 / 0.7995	1.0427 / 1.0223	0.4639 / 0.6455	0.7569 / 0.6178	0.7032 / 0.5664

Figure 57: Values of metrics for embeddings acquired by UMAP, second layer.

	ds	dd	cs	cd	acc_knn
filter 1	0.1880 / 0.3332	1.0903 / 1.0741	0.1350 / 0.2420	1.0515 / 0.9892	0.8804 / 0.7640

Figure 58: Values of metrics for embeddings acquired by UMAP, third layer.






Figure 59: Average value of the  $d_s$  metric (UMAP).




Figure 60: Average value of the  $d_d$  metric (UMAP).

## 6 Future work

We only used two very simple datasets, MNIST and FashionMNIST. Applying our methods to more datasets is needed to ensure our results are not specific to just these two datasets.

The metrics of embedding quality we used were rather arbitrary. A further research into measures of embedding quality and their subsequent application might make the obtained results more relevant.

The results were acquired by training the neural network once. To make the results more reliable, several runs and a subsequent statistical analysis would be a great step forward in terms of the quality of the obtained results.